

Reframing preprocessing selection as model-internal calibration in near-infrared spectroscopy

A large-scale benchmark of operator-adaptive PLS and Ridge models

Grégory Beurier^{1,2} Robin Reiter^{1,2} Camille Noûs³
Lauriane Rouan^{1,2} Denis Cornet^{1,2}

¹CIRAD, UMR AGAP Institut, Montpellier, France

²UMR AGAP Institut, Univ Montpellier, CIRAD, INRAE, Institut Agro, Montpellier, France

³Laboratoire Cogitamus, Havre de recherche, <https://www.cogitamus.fr/>
Camille Noûs: camille.nous@cogitamus.fr

May 17, 2026

Abstract

Preprocessing screening is often the most expensive part of a near-infrared spectroscopy calibration workflow. It works because smoothing, derivatives, detrending and related filters change the spectral directions seen by PLS or Ridge regression, but a full external search repeatedly refits nearly the same linear model. This paper studies the case where that search can be collapsed into one calibration step.

For strict linear preprocessing operators, the transformed PLS cross-covariance satisfies $(\mathbf{X}\mathbf{A}^\top)^\top\mathbf{Y} = \mathbf{A}\mathbf{X}^\top\mathbf{Y}$, and Ridge regression depends on the operator-induced kernel $\mathbf{X}\mathbf{A}^\top\mathbf{A}\mathbf{X}^\top$. These identities allow a finite operator bank to be screened inside the model while retaining original-wavelength coefficients. Sample-adaptive or fitted corrections such as SNV, MSC, EMSC and ASLS remain fold-local branches, not absorbed into the algebra.

The study uses the AOM benchmark cohort: 61 regression rows and 17 classification rows in the manifest. On the main regression denominator ($N = 32$), plain compact-bank AOM-PLS records median RMSEP ratios of 0.991 against PLS-default and 0.990 against PLS-HPO; the selected ASLS-AOM-compact-cv5 branch records 0.985 and 1.002 on the same two references. The plain AOMRidge-global-compact-none baseline records 0.974 against Ridge-default and 0.984 against Ridge-HPO, while the selected AOMRidge-Blender-headline-spxy3 records 0.918 and 0.966. The selected classifier, AOM-PLS-DA-global-simpls-covariance, improves balanced accuracy by 0.159 on $N = 13$ datasets with 12/13 wins. The runtime gap is the practical result: PLS-HPO takes a median total time of 710.81 s per run, whereas the selected AOM-PLS branch takes 1.63 s. Linear operator-adaptive calibration therefore gives comparable prediction quality to exhaustive preprocessing screening, with orders-of-magnitude less fitting time for PLS.

Keywords: near-infrared spectroscopy; chemometrics; preprocessing selection; partial least squares; ridge regression; operator kernels; calibration.

1 Introduction

Near-infrared spectroscopy (NIRS) is used because it is rapid, non-destructive and compatible with routine analytical deployment [1, 2]. The calibration step is still fragile. A model that works on raw spectra may fail after a different split, instrument or sample family, while the same

regressor with a derivative, smoother or baseline correction can become usable. In practice, much of the model-building effort is therefore spent screening preprocessing choices.

The usual response is external search. A calibration workflow enumerates normalizations, smoothers, derivative orders, baseline corrections and model hyperparameters, then chooses the best validated pipeline. That workflow is straightforward and often effective, but it is slow. It also hides a scientific question inside a large engineering loop: which spectral linear geometry does the calibration model need?

The cost is concrete in this benchmark. The full PLS preprocessing grid evaluates 600 combinations across five trials, and the Ridge grid evaluates the same combinations across ten trials, giving 3000 and 6000 candidate fits per dataset/seed before the final refit. Repeating this search on small calibration sets can also make the selected pipeline unstable: the chosen preprocessing is the winner of a multiplicative loop over folds, operators and model hyperparameters. This is the model-selection-bias regime familiar from the machine-learning literature [3–5], where a large inner search consumes calibration variance and inflates apparent skill. A single operator-adaptive linear fit is easier to audit because the selected geometry is part of one calibration object with coefficients recovered on the original wavelength axis.

This paper focuses on the part of preprocessing that has a simple algebraic answer. If a preprocessing step is a fixed linear operator on the wavelength axis, it can be represented by a matrix $\mathbf{A}$ and moved inside PLS or Ridge calibration. The model can then screen a bank of such operators in one fit rather than refitting every transformed data matrix as an external pipeline. This does not claim that every preprocessing step is linear. SNV, MSC, EMSC and ASLS are fitted or sample-adaptive transformations; they remain fold-local branches in the experimental protocol.

The contribution is deliberately narrow. We derive the PLS covariance identity and the Ridge operator-kernel identity, instantiate them in AOM-PLS and AOM-Ridge, and test the resulting models against default PLS/Ridge and full cartesian preprocessing-HPO baselines. The question is not whether a new nonlinear learner can win a benchmark. The question is whether the preprocessing search already used with linear chemometric models can be made faster, simpler and easier to audit.

The key empirical check is intentionally plain. The AOM-compact-cv5 PLS row and the AOMRidge-global-compact-none Ridge row use compact strict-linear banks without the selected ASLS branch, Blender selector or local-neighbourhood machinery. If these simple AOM baselines already compete with default and preprocessing-HPO PLS/Ridge, the algebra is doing useful work rather than only moving complexity into another tuning grid.

2 Related work

Spectral preprocessing is a central part of chemometrics. Common tools include Savitzky–Golay smoothing and derivatives [6], Norris–Williams derivatives [7], SNV [8], MSC [9], EMSC [10] and penalized baseline smoothers [11, 12]. Reviews emphasize that these transformations remove different nuisance structures and should be chosen with the analytical setting in mind [13–15]. A recurring observation in this literature is that no single preprocessing wins across datasets, so any practical workflow has to choose among a family of candidates, often through an outer cross-validated search.

Several authors have proposed combining multiple preprocessed blocks rather than selecting one. Sequential Preprocessing through Orthogonalization (SPORT) [16], Parallel Orthogonalized Partial Least Squares (PORTO) [17] and PROSAC [18] use multiblock PLS to fuse several preprocessed views of the same spectra; a broader ensemble framing is given in Mishra et al. [19]. These methods keep preprocessing outside the model and stack the alternatives downstream, which is a complementary strategy to the one explored here.




Figure 1: External preprocessing screening repeatedly fits transformed pipelines. Operator-adaptive calibration moves strict linear alternatives inside the calibration model; fitted or sample-adaptive corrections remain fold-local branches.

PLS remains the reference linear method for many NIRS calibration problems [20, 21]. NIPALS and SIMPLS offer two useful views: one iterative and score-based, the other covariance-based [22]. Ridge regression provides a complementary regularized linear model [23]; in dual form it depends on the sample Gram matrix, which makes preprocessing a change of kernel geometry [24, 25].

This work combines these two lines. It keeps the chemometric interpretability of linear calibration, but, in contrast with SPORT, PORTO and PROSAC, it treats strict preprocessing alternatives as fixed linear operators that can be screened algebraically inside one PLS or Ridge calibration rather than fused across multiple separately fitted preprocessed blocks. The result is a single deployable linear model whose spectral coefficients live on the original wavelength grid.

3 Linear-operator scope

Let $\mathbf{X} \in \mathbb{R}^{n \times p}$ be a centered spectral matrix, with rows as samples and columns as wavelengths. A strict linear preprocessing operator is a fixed matrix $\mathbf{A} \in \mathbb{R}^{p \times p}$ applied to each row spectrum:

$$\mathbf{X}_{\mathbf{A}} = \mathbf{X}\mathbf{A}^{\top}. \quad (1)$$

The operator is fixed once the wavelength grid and operator parameters are chosen. It does not depend on the response, the validation fold, the current sample, or a reference spectrum estimated from the calibration set.

The strict-linear bank used in the main experiments contains identity, Savitzky–Golay smoothers, Savitzky–Golay first and second derivatives, finite differences, and polynomial detrending projections. The compact bank is intentionally small: it has nine operators, so internal CV selection evaluates a small number of candidate geometries rather than a large cartesian preprocessing grid.

The outside-scope examples are important. SNV uses sample-specific centering and scaling. MSC and EMSC estimate correction parameters relative to a reference. ASLS estimates a baseline through an iterative asymmetric criterion. These transformations are allowed in the

Table 1: Operator-bank scope. Strict linear operators are handled by the AOM identities; fitted or sample-adaptive transformations remain explicit fold-local branches.

Family	Operators in compact bank	Strict linear status
Identity	identity transform	yes
Smoothing	Savitzky–Golay smoothing, windows 11 and 21	yes
Derivatives	Savitzky–Golay first derivative, windows 11 and 21	yes
Derivatives	Savitzky–Golay second derivative, window 11	yes
Baseline trend	polynomial detrending, degrees 1 and 2	yes, fitted inside calibration folds
Local contrast	first finite difference	yes
Scatter correction	SNV, MSC, EMSC	branch preprocessing, fold-local
Baseline correction	ASLS and related asymmetric smoothers	branch preprocessing, fold-local

benchmark, but only as branches fitted inside each calibration fold. The selected PLS result uses such a branch followed by the compact strict-linear AOM bank.

4 Methods

4.1 Covariance identity for one-shot AOM-PLS

For centered $\mathbf{X}$ and $\mathbf{Y}$, the transformed PLS cross-covariance for operator $\mathbf{A}_b$ is

$$(\mathbf{X}\mathbf{A}_b^\top)^\top \mathbf{Y} = \mathbf{A}_b \mathbf{X}^\top \mathbf{Y}. \quad (2)$$

Equation 2 is the core simplification. It means that a bank of strict linear preprocessings can be screened by acting on $\mathbf{X}^\top \mathbf{Y}$, without fitting one full PLS model for every transformed matrix.

The simulation check for this identity is direct: explicit transformed-space PLS and covariance-space AOM-PLS must select the same operator direction when the same centered fold and operator bank are supplied. The check is used as an implementation guard, while the benchmark below evaluates the effect on held-out prediction.

4.2 AOM-PLS solvers and coefficient recovery

A standard PLS coefficient matrix can be written as

$$\mathbf{B} = \mathbf{W}(\mathbf{P}^\top \mathbf{W})^+ \mathbf{Q}^\top, \quad (3)$$

where $\mathbf{W}$ are weights, $\mathbf{P}$ are spectral loadings, $\mathbf{Q}$ are response loadings and $+$ denotes the Moore–Penrose inverse. In AOM-PLS, a component may be selected in a transformed space. If the transformed-space direction is r_a under operator $\mathbf{A}_b$, its original-wavelength direction is

$$z_a = \mathbf{A}_b^\top r_a. \quad (4)$$

With $\mathbf{Z} = [z_1, \dots, z_k]$ and scores $\mathbf{T} = \mathbf{X}\mathbf{Z}$, the recovered coefficient matrix is

$$\mathbf{B} = \mathbf{Z}(\mathbf{P}^\top \mathbf{Z})^+ \mathbf{Q}^\top. \quad (5)$$

The fitted model is therefore still one linear calibration on the original wavelength grid.

Two solver views are implemented. The NIPALS-adjoint solver applies operator adjoints during iterative score extraction. The SIMPLS-covariance solver uses Equation 2 directly in

covariance space. The main plain baseline is **AOM-compact-cv5**: the compact nine-operator bank with internal five-fold selection and no fitted ASLS branch. The selected regression result uses the same compact CV5 branch after ASLS is fitted fold-locally. The default bank, SIMPLS-only variants, and other non-selected AOM variants are reported in the supplementary material.

The compact bank was chosen from the selector diagnostics, not by adding an unbounded operator catalogue. In the refreshed diagnostics, the three most frequent strict operators (**sg_smooth_w21_p3**, **sg_d1_w21_p3** and **sg_d1_w11_p2**) account for 54.7% of compact-bank component selections, while identity remains selected on 8.7% of components. The operator bank is therefore diverse enough to adapt, but small enough to keep the selection problem stable.

4.3 AOM-Ridge operator kernels

For centered spectra $\mathbf{X}_c$ and response $\mathbf{Y}_c$, Ridge regression can be solved in the dual. Operator $\mathbf{A}_b$ induces the kernel

$$\mathbf{K}_b = \mathbf{X}_c \mathbf{A}_b^\top \mathbf{A}_b \mathbf{X}_c^\top. \quad (6)$$

For regularization $\alpha > 0$,

$$\mathbf{C} = (\mathbf{K}_b + \alpha \mathbf{I}_n)^{-1} \mathbf{Y}_c, \quad \beta_b = \mathbf{A}_b^\top \mathbf{A}_b \mathbf{X}_c^\top \mathbf{C}. \quad (7)$$

Weighted operator mixtures are obtained by replacing $\mathbf{K}_b$ with

$$\mathbf{K} = \sum_b s_b^2 \mathbf{X}_c \mathbf{A}_b^\top \mathbf{A}_b \mathbf{X}_c^\top, \quad (8)$$

which gives an original-space coefficient matrix through the matching operator metric $\sum_b s_b^2 \mathbf{A}_b^\top \mathbf{A}_b$.

The plain Ridge baseline is **AOMRidge-global-compact-none**: one global selection over the compact strict-linear bank, without Blender mixing or local neighbourhood routing. The selected Ridge result uses the Blender selector (**AOMRidge-Blender-headline-spxy3**). It combines several strong operator geometries through internal validation rather than committing to one noisy winner. GCV, local-neighbourhood and additional global compact selectors are mathematically described in the supplement.

5 Datasets and protocol

5.1 Simulation checks

The implementation checks use explicit-vs-operator simulations for the PLS covariance identity and the Ridge kernel identity. These checks verify that for fixed folds and fixed strict operators, the folded computation matches the explicit transformed computation used by an external search.

5.2 Datasets

The benchmark cohort is intentionally heterogeneous. It combines public and local NIRS calibration tasks from leaf physiology, fruit quality, grain and seed traits, dairy, beverages, meat quality, petroleum, soil amendments, wood products and pharmaceutical tablets. Several source databases contribute more than one response variable, so a row denotes a concrete prediction or classification task rather than only a spectral database. The cohort covers small calibration settings, high-sample leaf datasets and several high-dimensional low- n regimes.




Figure 2: AOM mathematical structure. AOM-PLS screens operators through cross-covariances; AOM-Ridge screens operator-induced kernels. Both recover coefficients on the original wavelength grid.

Table 2: Dataset shape summary for the AOM cohort. Here n is the total calibration plus external-test sample count, p is the number of spectral variables, C is the number of classes and I is the largest-class share for classification rows.

Task	N	n_{median}	n_{min}	n_{max}	p_{median}	p_{min}	p_{max}	$(p/n)_{\text{median}}$	C_{median}	I_{median}
Classification	17	511	56	7323	1951	235	2177	2.055	2	0.513
Regression	61	402	40	45417	1023	125	4200	2.382	—	—

Table 2 summarizes the cohort in the same compact form used for the companion benchmark inventory. Regression contains 61 rows with median $n = 402$ samples and median $p = 1023$ spectral variables. The classification manifest contains 17 rows, of which 16 are included in the benchmark analyses after one missing-file exclusion. Across the 17 manifest rows, the median n is 511 and the median p is 1951; the median number of classes is two and the median largest-class share is 0.513. The complete per-dataset manifest description is provided in the supplementary longtable, while all score tables below use their explicitly stated paired denominators.

Whenever a source dataset already defined an external split, that split was preserved. Otherwise deterministic SPXY was used for regression, with the stratified variant used for classification when class labels were present [26, 27]. Test sets were kept outside preprocessing, operator, component and regularization selection. No explicit outlier removal was applied during cohort construction: calibration and test samples were kept as supplied by the source datasets.

5.3 Benchmark protocol

The empirical benchmark reports eight regression variants and one selected classification variant. All paired comparisons reported below are computed on the strict intersection $N_{\cap} = 32$ of cohort rows available for every paper variant, so that ratios and win counts are directly comparable across variant pairs. The classification cohort contains 17 manifest rows, with 13 paired datasets for the selected classifier.

All selection steps are internal to the calibration set. RMSEP is the primary regression metric. Because response scales differ across traits, aggregate regression results use paired




Figure 3: Dataset diversity in the AOM cohort. Points show total sample count and number of spectral variables on logarithmic axes; colours separate regression and classification rows, and the inset summarizes coarse application domains.

RMSEP ratios. Classification uses balanced accuracy. Paired tests use one-sided Wilcoxon signed-rank tests with Holm correction within the displayed family.

The cartesian-HPO baselines are intentionally strong and expensive. They search the conventional NIRS preprocessing grid $norm \times smooth \times baseline \times osc \times n_{components}$ where $norm \in \{\text{none}, \text{SNV}, \text{MSC}, \text{EMSC}_2\}$, $smooth$ ranges over Savitzky–Golay smoothers and derivatives at multiple window/polynomial-order settings plus Gaussian smoothers, $baseline \in \{\text{none}, \text{detrnd}, \text{ASLS}\}$, $osc \in \{\text{none}, osc_1, osc_2, osc_3\}$ and $n_{components}$ varies through the range relevant for each dataset. This is the same space a chemometrician would screen by hand when assembling a strong conventional baseline (PLS or Ridge after SNV/MSC, Savitzky–Golay smoothing/derivatives and ASLS or detrnd baseline correction); the most frequently selected configuration across datasets behaves as a de-facto fixed-recipe baseline. The grid contains 600 preprocessing combinations. PLS-HPO drives the search with TabPFN-guided proposals at five trials per combination and Ridge-HPO at ten, for a total of 3000 and 6000 candidate fits per dataset/seed before the final refit.

The eight paper labels correspond to the following runner keys: PLS-default (`pls-default-cv5`), PLS-HPO (`pls-tabPFN-hpo-25trials`), AOM-PLS (simple) (`AOM-compact-cv5`), AOM-PLS (best) (`ASLS-AOM-compact-cv5`), Ridge-default (`ridge-default-cv5`), Ridge-HPO (`ridge-tabPFN-hpo-60trials`), AOM-Ridge (simple) (`AOMRidge-global-compact-none`) and AOM-Ridge (best) (`AOMRidge-Blender-headline-spxy3`). These runner keys are kept visible to bridge artifacts and tables; a paper-facing mapping table is given in the supplement.

6 Results

6.1 PLS regression

The plain AOM-PLS regression baseline is `AOM-compact-cv5`. Against PLS-default, it reaches a median RMSEP ratio of 0.991 on $N = 32$ datasets with 22/32 wins. Against PLS-HPO, it is

Table 3: Cohort denominators and data-quality rules used in the main manuscript.

Property	Summary used in this work
Regression manifest	61 included regression rows across the benchmark inventory.
Classification manifest	17 classification manifest rows; 16 included rows after one missing-file exclusion.
Main regression denominator	$N_{\cap} = 32$ datasets for the eight paper variants.
Paper variants	PLS-default, PLS-HPO, AOM-PLS (simple), AOM-PLS (best), Ridge-default, Ridge-HPO, AOM-Ridge (simple) and AOM-Ridge (best).
Analytical domains	Leaf physiology, fruit quality, grain and seed traits, dairy, beverages, meat quality, petroleum, soil, manure, wood products, tablets and public calibration datasets.
Sample and wavelength range	Training sets span 28–39,225 samples and 125–4,200 spectral variables in the local regression manifest.
Validation rule	External test sets are never used for preprocessing, operator, component or regularization selection.

near parity, with a median ratio of 0.990 on $N = 32$ and 19/32 wins. This row has no ASLS branch and no enlarged operator search; it is the compact strict-linear bank selected by internal five-fold validation.

The selected AOM-PLS regression variant is **ASLS-AOM-compact-cv5**. It is compared in the main text with the pre-registered linear references. Against PLS-default, it reaches a median RMSEP ratio of 0.985 on $N = 32$ datasets with 20/32 wins. Against PLS-HPO, it is essentially tied: median ratio 1.002 on $N = 32$, 15/32 wins.

This is the intended interpretation. Full external preprocessing-HPO can recover the benefit of careful preprocessing, but the gain is obtained by repeatedly fitting the linear calibration under a much larger search budget. AOM-PLS makes the strict-linear part of the same decision inside one coefficient-bearing calibration object.

6.2 Ridge regression

The plain Ridge AOM baseline is **AOMRidge-global-compact-none**. It improves over default Ridge with a median RMSEP ratio of 0.974 on $N = 32$ datasets, 25/32 wins and Holm-adjusted $p = 0.007$. Against Ridge-HPO, it remains close to the grid baseline: median ratio 0.984 on $N = 32$, 19/32 wins.

The selected Ridge variant is **AOMRidge-Blender-headline-spxy3**. It improves strongly over default Ridge: median RMSEP ratio 0.918 on $N = 32$ datasets with 27/32 wins and Holm-adjusted $p = 2.6 \times 10^{-4}$. It also improves over Ridge-HPO, with median ratio 0.966 on $N = 32$, 25/32 wins and $p = 0.033$.

The Ridge result supports the operator-kernel view. In Ridge, preprocessing changes the regularized sample geometry. Blender is the best main selector because it combines strong operator geometries through internal validation while preserving a single deployable Ridge model.

6.3 Classification

The selected classification variant is **AOM-PLS-DA-global-simpls-covariance**. On the paired classification cohort, it improves balanced accuracy by a median 0.159 on $N = 13$ datasets, with 12/13 wins and Holm-adjusted $p = 0.007$. This secondary result confirms that the covariance-space operator selection used for regression transfers to PLS-DA when the response encoding and scoring metric are changed.

Table 4: Main paired regression results. Ratios below one favour the row method over the reference named in the comparison.

Comparison	Evidence source	N	Median RMSEP ratio / wins; p_{Holm}
AOM-PLS (simple) vs PLS-default	AOM-PLS seeds012 / default-CV all	32	0.991; 22/32; 0.896
AOM-PLS (best) vs PLS-default	AOM-PLS seeds012 / default-CV all	32	0.985; 20/32; 1.000
AOM-PLS (simple) vs PLS-HPO	AOM-PLS seeds012 / cartesian HPO	32	0.990; 19/32; 1.000
AOM-PLS (best) vs PLS-HPO	AOM-PLS seeds012 / cartesian HPO	32	1.002; 15/32; 1.000
PLS-HPO vs PLS-default	cartesian HPO seeds012 / default-CV all	32	0.992; 19/32; 1.000
Ridge-HPO vs Ridge-default	cartesian HPO seeds012 / default-CV all	32	0.962; 19/32; 1.000
AOM-Ridge (simple) vs Ridge-default	AOM-Ridge headline / default-CV all	32	0.974; 25/32; 0.007
AOM-Ridge (best) vs Ridge-default	AOM-Ridge headline / default-CV all	32	0.918; 27/32; 2.6e-04
AOM-Ridge (simple) vs Ridge-HPO	AOM-Ridge headline / cartesian HPO seeds012	32	0.984; 19/32; 1.000
AOM-Ridge (best) vs Ridge-HPO	AOM-Ridge headline / cartesian HPO seeds012	32	0.966; 25/32; 0.033

Table 5: Main classification result. Positive balanced-accuracy differences favour the operator-adaptive classifier.

Comparison	N	Median Δ balanced acc.	95% CI	Wins; p_{Holm}
AOM-PLS-DA-global-simpls-covariance vs PLS-DA	13	0.159	0.129–0.422	12/13; 0.007

6.4 Time budget

The practical advantage is runtime. PLS-HPO has a median total time of 710.81 s per run and an observed row total of 40.8 h. AOM-PLS (simple), implemented as `AOM-compact-cv5`, has a median total time of 1.18 s, and AOM-PLS (best) has a median total time of 1.63 s and an observed total of 0.1 h. On median total time, AOM-PLS (best) is about 436 times faster than PLS-HPO. In runner-level search units, the contrast is 3000 trial fits per PLS-HPO dataset/seed versus 45 compact-bank CV fits for each AOM-PLS row.

Ridge is different because AOM-Ridge (best), the Blender selector, is heavier. It still improves over Ridge-HPO, and its median total time is 728.81 s compared with 1584.00 s for Ridge-HPO. AOM-Ridge (simple), implemented as `AOMRidge-global-compact-none`, is much lighter, with median total time 23.78 s and 450 operator-alpha cells. The result is not that every AOM selector is instant; it is that strict-linear preprocessing can be represented as a model geometry rather than as a blind external pipeline loop.




Figure 4: Eight paired regression comparisons for the plain compact-bank AOM baselines and the selected AOM-PLS and AOM-Ridge variants against their PLS-default, PLS-HPO, Ridge-default and Ridge-HPO references.

6.5 Failure modes

The headline ratios summarize medians; they hide a small number of datasets where AOM does not win. Two manifest rows (`FinalScore_grp70_30_scoreQ` and `Tleaf_grp70_30`) contain NaN values in the source spectra and fail under every linear regressor, including the PLS and Ridge baselines, so they are excluded from paired comparisons by construction rather than by AOM failure. Among the rows where RMSEPs are finite, AOM-PLS is weaker than the PLS-HPO winner on the `Brix_spxy70` berry split, on the `Rice_Amylose_313_YbasedSplit` amylose row, on the `colza C_woOutlier` row and on the `brix_groupSampleID_stratDateVar_balRows` strawberry-puree row. These are datasets where SNV, MSC or a fitted baseline correction carries most of the calibration value and the strict linear bank captures only part of it; they are also the rows that drive the upper tail of the paired ratio distribution for `ASLS-AOM-compact-cv5`. Three further datasets (`LMA_spxyG70_30_byCultivar_ASD`, `N_woOutlier`, `N_woOutlier`) are included in the supplementary failure-modes table not because the selected AOM-PLS variant loses on them — it does not — but because the POP-PLS family records substantially higher RMSEP than the covariance-space AOM-PLS rows on the same datasets, which is the diagnostic evidence that motivates dropping POP-PLS from the headline. The supplementary *Failure modes* table reports the finite-RMSEP range across all nine logged AOM-PLS and POP-PLS variants for both classes of row, and the supplementary *Per-dataset prediction quality* section gives the per-dataset gain figure that visualizes the same picture.

7 Discussion

The benchmark supports a simple engineering conclusion: when preprocessing alternatives are strict linear operators, screening them outside the model is often unnecessary. AOM-PLS folds the operator choice into the PLS cross-covariance calculation, and AOM-Ridge folds it into the Ridge kernel. The resulting calibration remains linear and auditable.

The strongest use case is the one common in NIRS: small to moderate calibration sets, noisy validation folds and many plausible preprocessing choices. In that regime, an exhaustive external search spends most of its time refitting variations of the same linear model. The grid

Table 6: Search and runtime budget for the main methods. Observed totals are sums over benchmark rows.

Method	Datasets	Search budget	Median fit (s)	Median total (s)
PLS-default	32	25 component trials; observed 0.1 h	0.02	1.21
PLS-HPO	32	600 x 5 trials; ob- served 40.8 h	0.03	710.81
AOM-PLS (simple)	32	9 operators x CV-5; observed 0.1 h	1.18	1.18
AOM-PLS (best)	32	ASLS branch + 9 operators x CV-5; observed 0.1 h	1.43	1.63
Ridge-default	32	15 alpha trials; ob- served 0.0 h	0.00	0.38
Ridge-HPO	32	600 x 10 trials; ob- served 96.5 h	0.05	1584.00
AOM-Ridge (simple)	32	9 operators x 50 al- pha cells; observed 0.5 h	23.77	23.78
AOM-Ridge (best)	32	AOM-Ridge Blender: 8 can- didates x 3 outer folds + 8 refits; observed 9.3 h	727.51	728.81

baselines here require 3000 PLS or 6000 Ridge candidate fits per dataset/seed, whereas the compact AOM-PLS selector evaluates 45 operator/component candidates. The observed median total times reflect that difference: 710.81 s for PLS-HPO against 1.18–1.63 s for the compact AOM-PLS rows.

The statistical reason is related to the engineering cost. A large cartesian grid repeatedly re-estimates pipelines on small folds and then promotes the winner of that noisy search. The strict-linear AOM step instead selects operators from a covariance or kernel representation of one calibration problem. It does not remove validation uncertainty, but it avoids turning the choice of a derivative or smoother into a multiplicative refitting loop.

Auditability and re-deployment are also different. A grid winner is an external preprocessing pipeline plus a downstream linear model; every fitted branch must be reproduced correctly when the calibration is moved to another instrument or software stack. AOM-PLS and AOM-Ridge recover coefficients on the original wavelength axis (Equations 5 and 7), so the deployed object is still one linear calibration whose spectral coefficients can be inspected directly.

The plain baselines are the most useful evidence for choosing this route over a preprocessing grid. `AOM-compact-cv5` already improves modestly over default PLS and reaches near-parity with PLS-HPO, while `AOMRidge-global-compact-none` improves over default Ridge and remains close to Ridge-HPO. The selected ASLS-AOM and Blender rows add accuracy, but the compact no-frills rows show that the main gain is not dependent on a large new tuning apparatus.

The scope is also clear. AOM algebra does not absorb sample-adaptive transforms. SNV, MSC, EMSC and ASLS must be fitted inside folds. This is why the selected PLS result is described as an ASLS branch followed by strict linear AOM, not as ASLS being part of the operator identity. When such branches dominate a dataset, they should remain explicit and audited.

Taken together, the simple-baseline result is the practical summary: before adding the selected ASLS or Blender refinements, compact AOM already captures much of the value of

Table 7: Software artifacts and validation status.

Component	Tests / evidence	Status
nirs4all-aom Python package	AOM-PLS, POP-PLS, AOM-Ridge and FastAOM implementations, with unit tests and sklearn-compatible APIs.	public software
nirs4all-aom benchmark artifacts	Benchmark runners, result CSVs, aggregation scripts, figure builders and manuscript tables used here.	release with paper
nirs4all instrumentation context	NIRS instrumentation, acquisition and provenance context for local benchmark inputs.	instrumentation software
Supplementary validation dossier	Cohort manifest, missing-dataset audit, paired statistics and software-readiness notes distributed with nirs4all-aom .	public evidence

preprocessing search while keeping the model linear, fast and directly deployable.

8 Conclusion

Preprocessing screening for linear NIRS calibration works, but it is often implemented as an expensive outer loop. For strict linear operators, PLS and Ridge provide identities that move the operator choice inside the calibration itself. On the available AOM benchmark, the selected AOM-PLS and AOM-Ridge variants are comparable to or better than full preprocessing-HPO on the reported paired tests, with a large time advantage for PLS. The useful scientific message is therefore modest and practical: keep fitted corrections fold-local, fold strict linear operators into the model, and report every comparison on its paired denominator.

Data and code availability

The public software entry point for the methods, benchmark runners, result tables and manuscript artifacts reported here is **nirs4all-aom** [28] at <https://github.com/GBeurier/nirs4all-aom>. The **nirs4all** project [29] is cited only for the NIRS instrumentation, acquisition and provenance context used to assemble local benchmark inputs; it is not the software entry point for the AOM methods reported in this paper.

The reproducibility protocol is defined by the cohort denominators in Section 5, the fold-local model-selection rules in the Methods, and paired dataset-level tests with Holm correction for the reported families of comparisons. The supplementary material gives the full cohort manifest, per-dataset regression table, prediction-quality diagnostics, seed-stability diagnostics and software-validation matrix. Table 7 summarizes the software artifacts and validation status for the public release.

Supplementary Material

The supplementary material contains complete derivations, the cohort manifest, operator-bank diagnostics, non-selected AOM variants, alternative AOM-form explorations, per-dataset tables, classification details, seed stability, software artifacts and the AI-assistance statement.

Declaration of Generative AI and AI-assisted technologies in the writing process

During the preparation of this work, the authors used Anthropic Claude and OpenAI Codex to assist with code review, implementation, repository organization, benchmark aggregation scripts, LaTeX editing and drafting or revision of explanatory text. After using these tools, the authors reviewed, edited and verified the code, numerical results, references and manuscript content as needed, and take full responsibility for the content of the publication.

References

- [1] Donald A. Burns and Emil W. Ciurczak. *Handbook of Near-Infrared Analysis*. CRC Press, 3 edition, 2007.
- [2] Celio Pasquini. Near infrared spectroscopy: A mature analytical technique with new perspectives. *Analytica Chimica Acta*, 1026:8–36, 2018. doi: 10.1016/j.aca.2018.04.004.
- [3] Sudhir Varma and Richard Simon. Bias in error estimation when using cross-validation for model selection. *BMC Bioinformatics*, 7:91, 2006. doi: 10.1186/1471-2105-7-91.
- [4] Gavin C. Cawley and Nicola L. C. Talbot. On over-fitting in model selection and subsequent selection bias in performance evaluation. *Journal of Machine Learning Research*, 11:2079–2107, 2010.
- [5] James Bergstra and Yoshua Bengio. Random search for hyper-parameter optimization. *Journal of Machine Learning Research*, 13:281–305, 2012.
- [6] Abraham Savitzky and Marcel J. E. Golay. Smoothing and differentiation of data by simplified least squares procedures. *Analytical Chemistry*, 36(8):1627–1639, 1964. doi: 10.1021/ac60214a047.
- [7] Karl H. Norris and Phil C. Williams. Influence of moisture content on the reflective behavior of grain. *Cereal Chemistry*, 53(6):794–805, 1976.
- [8] R. J. Barnes, M. S. Dhanoa, and Susan J. Lister. Standard normal variate transformation and de-trending of near-infrared diffuse reflectance spectra. *Applied Spectroscopy*, 43(5):772–777, 1989.
- [9] Paul Geladi, Donald MacDougall, and Harald Martens. Linearization and scatter-correction for near-infrared reflectance spectra of meat. *Applied Spectroscopy*, 39(3):491–500, 1985. doi: 10.1366/0003702854248656.
- [10] Harald Martens and Edward Stark. Extended multiplicative signal correction and spectral interference subtraction. *Journal of Pharmaceutical and Biomedical Analysis*, 9(8):625–635, 1991.
- [11] Paul H. C. Eilers. A perfect smoother. *Analytical Chemistry*, 75(14):3631–3636, 2003. doi: 10.1021/ac034173t.
- [12] Zhi-Min Zhang, Shan Chen, and Yi-Zeng Liang. Baseline correction using adaptive iteratively reweighted penalized least squares. *Analyst*, 135(5):1138–1146, 2010.
- [13] Åsmund Rinnan, Frans van den Berg, and Søren Balling Engelsen. Review of the most common pre-processing techniques for near-infrared spectra. *TrAC Trends in Analytical Chemistry*, 28(10):1201–1222, 2009. doi: 10.1016/j.trac.2009.07.007.

- [14] Jasper Engel, Jan Gerretzen, Ewa Szymańska, Jeroen J. Jansen, Gerard Downey, Lionel Blanchet, and Lutgarde M. C. Buydens. Breaking with trends in pre-processing? *TrAC Trends in Analytical Chemistry*, 50:96–106, 2013. doi: 10.1016/j.trac.2013.04.015.
- [15] Jean-Michel Roger, Alexia Mallet, and Federico Marini. Preprocessing NIR spectra for aquaphotomics. *Molecules*, 27(20):6795, 2022. doi: 10.3390/molecules27206795.
- [16] Jean-Michel Roger, Alessandra Biancolillo, and Federico Marini. SPORT pre-processing can improve near-infrared quality prediction models for fresh fruits and agro-materials. *Chemometrics and Intelligent Laboratory Systems*, 199:103975, 2020. doi: 10.1016/j.chemolab.2020.103975.
- [17] Puneet Mishra, Jean-Michel Roger, Douglas N. Rutledge, and Federico Marini. A brief note on the PORTO algorithm for multi-block analysis with applications to spectral imaging. *Chemometrics and Intelligent Laboratory Systems*, 212:104190, 2021. doi: 10.1016/j.chemolab.2021.104190.
- [18] Puneet Mishra, Jean-Michel Roger, Federico Marini, Alessandra Biancolillo, and Douglas N. Rutledge. Pre-processing ensembles with response oriented sequential alternation calibration (PROSAC): A step towards ending the pre-processing search and optimization quest for near-infrared spectral modelling. *Chemometrics and Intelligent Laboratory Systems*, 222:104497, 2022. doi: 10.1016/j.chemolab.2022.104497.
- [19] Puneet Mishra, Alessandra Biancolillo, Jean-Michel Roger, Federico Marini, and Douglas N. Rutledge. New data preprocessing trends based on ensemble of multiple preprocessing techniques. *TrAC Trends in Analytical Chemistry*, 132:116045, 2020. doi: 10.1016/j.trac.2020.116045.
- [20] Paul Geladi and Bruce R. Kowalski. Partial least-squares regression: a tutorial. *Analytica Chimica Acta*, 185:1–17, 1986. doi: 10.1016/0003-2670(86)80028-9.
- [21] Svante Wold, Michael Sjöström, and Lennart Eriksson. Pls-regression: a basic tool of chemometrics. *Chemometrics and Intelligent Laboratory Systems*, 58(2):109–130, 2001. doi: 10.1016/S0169-7439(01)00155-1.
- [22] Sijmen De Jong. Simpls: an alternative approach to partial least squares regression. *Chemometrics and Intelligent Laboratory Systems*, 18(3):251–263, 1993.
- [23] Arthur E. Hoerl and Robert W. Kennard. Ridge regression: biased estimation for nonorthogonal problems. *Technometrics*, 12(1):55–67, 1970. doi: 10.1080/00401706.1970.10488634.
- [24] Bernhard Schölkopf and Alexander J. Smola. *Learning with Kernels*. MIT Press, 2002.
- [25] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning*. Springer, 2 edition, 2009. doi: 10.1007/978-0-387-84858-7.
- [26] R. W. Kennard and L. A. Stone. Computer aided design of experiments. *Technometrics*, 11(1):137–148, 1969.
- [27] Roberto K. H. Galvão, Mario C. U. Araujo, Gledson E. Jose, Marcio J. C. Pontes, Edvan C. Silva, and Teresa C. B. Saldanha. A method for calibration and validation subset partitioning. *Talanta*, 67(4):736–740, 2005. doi: 10.1016/j.talanta.2005.03.025.
- [28] Grégory Beurier. nirs4all-aom: Adaptive operator-mixture pls and ridge for nir spectroscopy. <https://github.com/GBeurier/nirs4all-aom>, 2026. Software, benchmark artifacts and manuscript sources.

- [29] Grégory Beurier, Denis Cornet, and Lauriane Rouan. nirs4all: Nirs instrumentation and acquisition toolkit. <https://github.com/GBeurier/nirs4all>, 2026. Cited for instrumentation, acquisition and benchmark provenance context.




Figure 5: Prediction-quality scatter plots for the selected AOM-PLS and AOM-Ridge variants against default and HPO references. Axes are log RMSEP; points below the diagonal favour AOM.




Figure 6: R^2 coverage curves for the eight paper variants on the main regression denominator. Higher curves indicate a larger fraction of datasets exceeding each test- R^2 threshold.




Figure 7: Search-budget scale for the eight paper variants: PLS-default, PLS-HPO, AOM-PLS (simple), AOM-PLS (best), Ridge-default, Ridge-HPO, AOM-Ridge (simple) and AOM-Ridge (best). Counts are runner-level search units: logged trials for PLS/Ridge defaults and HPO, compact-bank CV fits for AOM-PLS, operator-alpha cells for AOM-Ridge (simple), and top-level outer-selector fits plus refits for AOM-Ridge (best).




Figure 8: Accuracy/time view for all eight paper variants. The horizontal axis uses median total fit/search time: default and HPO rows include the complete CV or search/refit run, while AOM rows include the selector fit and final fitted model.




Figure 9: Runtime distributions for all eight paper variants. PLS-default, PLS-HPO, Ridge-default and Ridge-HPO boxes show total CV or HPO run time; AOM-PLS and AOM-Ridge boxes show selector fit time plus prediction overhead.